

Assignment 9: GBDT

Response Coding: Example

Train Data		Encoded Train Data		
State	class	State_0	State_1	class
A	0	3/5	2/5	0
B	1	0/2	2/2	1
C	1	1/3	2/3	1
A	0	3/5	2/5	0
A	1	3/5	2/5	1
B	1	0/2	2/2	1
A	0	3/5	2/5	0
A	1	3/5	2/5	1
C	1	1/3	2/3	1
C	0	1/3	2/3	0

Resonse table(only from train)			
State	Class=0	Class=1	
A	3	2	
B	0	2	
C	1	2	

Test Data		Encoded Test Data	
State		State_0	State_1
A		3/5	2/5
C		1/3	2/3
D		1/2	1/2
C		1/3	2/3
B		0/2	2/2
E		1/2	1/2

The response tabel is built only on train dataset. For a category which is not there in train data and present in test data, we will encode them with default values Ex: in our test data if have State: D then we encode it as [0.5, 0.05]

1. Apply GBDT on these feature sets


- **Set 1:** categorical(instead of one hot encoding, try [response coding](#): use probability values), numerical features + project_title(TFIDF)+ preprocessed_eassay (TFIDF)+sentiment Score of eassay(check the bellow example, include all 4 values as 4 features)
- **Set 2:** categorical(instead of one hot encoding, try [response coding](#): use probability values), numerical features + project_title(TFIDF W2V)+ preprocessed_eassay (TFIDF W2V)


2. The hyper paramter tuning (Consider any two hyper parameters)

- Find the best hyper parameter which will give the maximum [AUC](#) value
- find the best hyper paramter using k-fold cross validation/simple cross validation data
- use gridsearch cv or randomsearch cv or you can write your own for loops to do this task

3. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure






with X-axis as **n_estimators**, Y-axis as **max_depth**, and Z-axis as **AUC Score**, we have given the notebook which explains how to plot this 3d plot, you can find it in the same drive [3d_scatter_plot.ipynb](#)


or

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure



[seaborn heat maps](#) with rows as **n_estimators**, columns as **max_depth**, and values inside the cell representing **AUC Score**

- You choose either of the plotting techniques out of 3d plot or heat map
- Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.



- Along with plotting ROC curve, you need to print the [confusion matrix](#) with predicted and original labels of test data points

	Predicted: NO	Predicted: YES
Actual: NO	TN = ??	FP = ??
Actual: YES	FN = ??	TP = ??

- You need to summarize the results at the end of the notebook, summarize it in the table format

Vectorizer	Model	Hyper parameter	AUC
BOW	Brute	7	0.78
TFIDF	Brute	12	0.79

W2V	Brute	10	0.78
TfIDFW2V	Brute	6	0.78

1. GBDT (xgboost/lightgbm)

In [52]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
import re
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle
from tqdm import tqdm
import os
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter
```

In [170]:

```
with open(r'E:\assignment\donar case\glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

1.1 Loading Data

In [171]:

```
prep_data = pd.read_csv(r'E:\assignment\donar case\preprocessed_data.csv')
prep_data.head(2)
```

Out[171]:

	school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_posted_projects	project_is_approved	clean_category
0	ca	mrs	grades_prek_2	53	1	math_s
1	ut	ms	grades_3_5	4	1	special

1.2 Splitting data into Train and cross validation(or test): Stratified Sampling

In [172]:

```
import nltk
from nltk.sentiment.vader import SentimentIntensityAnalyzer
analyzer = SentimentIntensityAnalyzer()
neg=[];pos=[];neu=[]; compound = []
for i in range(len(prepare_data['essay'])):
    sentiment_scores = analyzer.polarity_scores(prepare_data['essay'][i])
    neg.append(sentiment_scores['neg'])
    pos.append(sentiment_scores['pos'])
    neu.append(sentiment_scores['neu'])
    compound.append(sentiment_scores['compound'])
```

In [173]:

```
#new columns indicating the sentiment score of each project essay
prepare_data['neg'] = neg
prepare_data['neu'] = neu
prepare_data['pos'] = pos
prepare_data['compound'] = compound
```

In [174]:

```
# train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(prepare_data, prepare_data['project_is_approved'], test_size=0.33, stratify=prepare_data['project_is_approved'])
```

1.3 Make Data Model Ready: encoding essay, and project_title

In [175]:

```
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer_tfidf_essay = TfidfVectorizer(min_df=10)
vectorizer_tfidf_essay.fit(X_train['essay']) #Fitting has to be on Train data
train_essay_tfidf = vectorizer_tfidf_essay.transform(X_train['essay'].values)
test_essay_tfidf =vectorizer_tfidf_essay.transform(X_test['essay'].values)
print("Shape of train data matrix after one hot encoding ",train_essay_tfidf.shape)
print("Shape of test data matrix after one hot encoding ",test_essay_tfidf.shape)
```

Shape of train data matrix after one hot encoding (73196, 14184)
Shape of test data matrix after one hot encoding (36052, 14184)

In [176]:

```
# average Word2Vec
# compute average word2vec for each review.
train_avg_w2v_essays = []; # the avg-w2v for each sentence/review is stored in this list
train_avg_w2v_essays = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in X_train['essay']: # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
        if cnt_words != 0:
            vector /= cnt_words
    train_avg_w2v_essays.append(vector)
print(len(train_avg_w2v_essays))
print(len(train_avg_w2v_essays[0]))
```

141
300

In [177]:

```
# average Word2Vec
# compute average word2vec for each review.
test_avg_w2v_essays = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in X_test['essay']: # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
        if cnt_words != 0:
            vector /= cnt_words
    test_avg_w2v_essays.append(vector)
print(len(test_avg_w2v_essays))
print(len(test_avg_w2v_essays[0]))
```

36052

300

In [178]:

```
# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
tfidf_model = TfidfVectorizer()
tfidf_model.fit(X_train['essay'].values)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
tfidf_words = set(tfidf_model.get_feature_names())
```

In [179]:

```
#average Word2Vec
# compute average word2vec for each review.
train_tfidf_w2v_essays = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in X_train['essay']: # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
# here we are multiplying idf value(dictionary[word]) and the tf
tfvalue((sentence.count(word)/len(sentence.split())))
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
idf value for each word
            vector += (vec * tf_idf) # calculating tfidf weighted w2vtf_idf_weight += tf_idf
        if tf_idf_weight != 0:
            vector /= tf_idf_weight
    train_tfidf_w2v_essays.append(vector)
print(len(train_tfidf_w2v_essays))
print(len(train_tfidf_w2v_essays[0]))
```

73196

300

In [180]:

```
#average Word2Vec
# compute average word2vec for each review.
test_tfidf_w2v_essays = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in X_test['essay']: # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
# here we are multiplying idf value(dictionary[word]) and the tf
value((sentence.count(word)/len(sentence.split())))
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
idf value for each word
            vector += (vec * tf_idf) # calculating tfidf weighted w2vtf_idf_weight += tf_idf
        if tf_idf_weight != 0:
            vector /= tf_idf_weight
```

```

test_tfidf_w2v_essays.append(vector)
print(len(test_tfidf_w2v_essays))
print(len(test_tfidf_w2v_essays[0]))

```

36052
300

1.4 Make Data Model Ready: encoding numerical, categorical features

In [181]:

#Response Coding

```

def mask(df, key, value):
    return df[df[key] == value]
def get_response(data, data_label):
    cat_values = np.unique(data).tolist()
    d=""
    d=data.name
    cat_test_values=np.unique(X_test[d]).tolist()
    df = pd.DataFrame({'feature':data.values.tolist(), 'label':data_label.values.tolist()})
    pd.DataFrame.mask = mask
    accep = {};reject={};prob_neg = {};prob_pos={}
    for i in cat_values:
        if i in cat_test_values:
            count_0 = len(df.mask('feature', i).mask('label', 0))
            count_1 = len(df.mask('feature', i).mask('label', 1))
            total = count_0 + count_1
            prob_0 = count_0/total
            prob_1 = count_1/total
            accep[i] = count_1
            reject[i] = count_0
            prob_neg[i] = prob_0
            prob_pos[i] = prob_1
        else:
            prob_neg[i] = 0.5
            prob_pos[i] = 0.5

    return prob_neg,prob_pos

```

In [182]:

```

cat_0_train = get_response(X_train['clean_categories'],y_train)[0]
cat_1_train = get_response(X_train['clean_categories'],y_train)[1]

```

In [183]:

```

subcat_0_train = get_response(X_train['clean_subcategories'],y_train)[0]
subcat_1_train = get_response(X_train['clean_subcategories'],y_train)[1]

```

In [184]:

```

state_0_train = get_response(X_train['school_state'],y_train)[0]
state_1_train = get_response(X_train['school_state'],y_train)[1]
prefix_0_train = get_response(X_train['teacher_prefix'],y_train)[0]
prefix_1_train = get_response(X_train['teacher_prefix'],y_train)[1]
grad_cat_0_train = get_response(X_train['project_grade_category'],y_train)[0]
grad_cat_1_train = get_response(X_train['project_grade_category'],y_train)[1]
cat_0_test = get_response(X_test['clean_categories'],y_test)[0]
cat_1_test = get_response(X_test['clean_categories'],y_test)[1]
subcat_0_test = get_response(X_test['clean_subcategories'],y_test)[0]
subcat_1_test = get_response(X_test['clean_subcategories'],y_test)[1]
state_0_test = get_response(X_test['school_state'],y_test)[0]
state_1_test = get_response(X_test['school_state'],y_test)[1]
prefix_0_test = get_response(X_test['teacher_prefix'],y_test)[0]
prefix_1_test = get_response(X_test['teacher_prefix'],y_test)[1]
grad_cat_0_test = get_response(X_test['project_grade_category'],y_test)[0]
grad_cat_1_test = get_response(X_test['project_grade_category'],y_test)[1]

```

In [185]:

```
cat_neg_train = []
cat_pos_train = []
for i in X_train['clean_categories']:
    cat_neg_train.append(cat_0_train[i])
    cat_pos_train.append(cat_1_train[i])
X_train['cat_0'] = cat_neg_train
X_train['cat_1'] = cat_pos_train
print(X_train['cat_0'].shape)
print(X_train['cat_1'].shape)
```

(73196,)

(73196,)

In [186]:

```
cat_neg_test = []
cat_pos_test = []
for i in X_test['clean_categories']:
    cat_neg_test.append(cat_0_test[i])
    cat_pos_test.append(cat_1_test[i])
X_test['cat_0'] = cat_neg_test
X_test['cat_1'] = cat_pos_test
print(X_test['cat_0'].shape)
print(X_test['cat_1'].shape)
```

(36052,)

(36052,)

In [187]:

```
subcat_neg_train = []
subcat_pos_train = []
for i in X_train['clean_subcategories']:
    subcat_neg_train.append(subcat_0_train[i])
    subcat_pos_train.append(subcat_1_train[i])
X_train['subcat_0'] = subcat_neg_train
X_train['subcat_1'] = subcat_pos_train
print(X_train['subcat_0'].shape)
print(X_train['subcat_1'].shape)
```

(73196,)

(73196,)

In [188]:

```
subcat_neg_test = []
subcat_pos_test = []
for i in X_test['clean_subcategories']:
    subcat_neg_test.append(subcat_0_test[i])
    subcat_pos_test.append(subcat_1_test[i])
X_test['subcat_0'] = subcat_neg_test
X_test['subcat_1'] = subcat_pos_test
print(X_test['subcat_0'].shape)
print(X_test['subcat_1'].shape)
```

(36052,)

(36052,)

In [189]:

```
state_neg_train = []
state_pos_train = []
for i in X_train['school_state']:
    state_neg_train.append(state_0_train[i])
    state_pos_train.append(state_1_train[i])
X_train['state_0'] = state_neg_train
X_train['state_1'] = state_pos_train
print(X_train['state_0'].shape)
print(X_train['state_1'].shape)
```

```
(73196,)
(73196,)
```

In [190]:

```
state_neg_test = []
state_pos_test = []
for i in X_test['school_state']:
    state_neg_test.append(state_0_test[i])
    state_pos_test.append(state_1_test[i])
X_test['state_0'] = state_neg_test
X_test['state_1'] = state_pos_test
print(X_test['state_0'].shape)
print(X_test['state_1'].shape)
```

```
(36052,)
(36052,)
```

In [191]:

```
prefix_neg_train = []
prefix_pos_train = []
for i in X_train['teacher_prefix']:
    prefix_neg_train.append(prefix_0_train[i])
    prefix_pos_train.append(prefix_1_train[i])
X_train['prefix_0'] = prefix_neg_train
X_train['prefix_1'] = prefix_pos_train
print(X_train['prefix_0'].shape)
print(X_train['prefix_1'].shape)
```

```
(73196,)
(73196,)
```

In [192]:

```
prefix_neg_test = []
prefix_pos_test = []
for i in X_test['teacher_prefix']:
    prefix_neg_test.append(prefix_0_test[i])
    prefix_pos_test.append(prefix_1_test[i])
X_test['prefix_0'] = prefix_neg_test
X_test['prefix_1'] = prefix_pos_test
print(X_test['prefix_0'].shape)
print(X_test['prefix_1'].shape)
```

```
(36052,)
(36052,)
```

In [193]:

```
grade_neg_train = []
grade_pos_train = []
for i in X_train['project_grade_category']:
    grade_neg_train.append(grad_cat_0_train[i])
    grade_pos_train.append(grad_cat_1_train[i])
X_train['grade_0'] = grade_neg_train
X_train['grade_1'] = grade_pos_train
print(X_train['grade_0'].shape)
print(X_train['grade_1'].shape)
```

```
(73196,)
(73196,)
```

In [194]:

```
grade_neg_test = []
grade_pos_test = []
for i in X_test['project_grade_category']:
```


(36052,)
(36052,)

#Concatenating all TFIDF Features

```
Final TFIDF Data matrix
(73196, 14232) (73196,)
(36052, 14232) (36052,)
```

#Applying XGBOOST ON TFIDF

[illegible]

```

interaction_constraints=None,
learning_rate=None, max_delta_step=None,
max_depth=None, min_child_weight=None,
missing=nan, monotone_constraints=None,
num_parallel_tree=None,
objective='binary:logistic',
random_state=None, reg_alpha=None,
reg_lambda=None, scale_pos_weight=None,
subsample=None, tree_method=None,
validate_parameters=False,
verbosity=None),

iid='warn', n_jobs=-1,
param_grid={'max_depth': [2, 5, 10],
            'n_estimators': [10, 50, 100]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring='roc_auc', verbose=0)

```

In [83]:

```

print('Best score: ',gs.best_score_)
print('k value with best score: ',gs.best_params_)
print('='*75)
print('Train AUC scores')
print(gs.cv_results_['mean_train_score'])
print('CV AUC scores')
print(gs.cv_results_['mean_test_score'])

```

```

Best score: 0.7261385332980901
k value with best score: {'max_depth': 2, 'n_estimators': 100}
=====
Train AUC scores
[0.6783479  0.74382595 0.77496039 0.74799896 0.85705737 0.90950417
 0.89988699 0.98307372 0.99671967]
CV AUC scores
[0.66984183 0.71499992 0.72613853 0.69371443 0.71995857 0.72169263
 0.69268183 0.71088139 0.70943098]


```

In [84]:

```

import seaborn as sns; sns.set()
max_scores1 = pd.DataFrame(gs.cv_results_).groupby(['param_n_estimators', 'param_max_depth']).max()
).unstack()['mean_test_score', 'mean_train_score']
fig, ax = plt.subplots(1,2, figsize=(20,6))
sns.heatmap(max_scores1.mean_train_score, annot = True, fmt='.4g', ax=ax[0])
sns.heatmap(max_scores1.mean_test_score, annot = True, fmt='.4g', ax=ax[1])
ax[0].set_title('Train Set')
ax[1].set_title('CV Set')
plt.show()

```



In [87]:

```
gs.best_params_
```

Out [87]:

```
{'max_depth': 2, 'n_estimators': 100}
```

In [86]:

```
max_d = gs.best_params_['max_depth']
n_est = gs.best_params_['n_estimators']
```

In [88]:

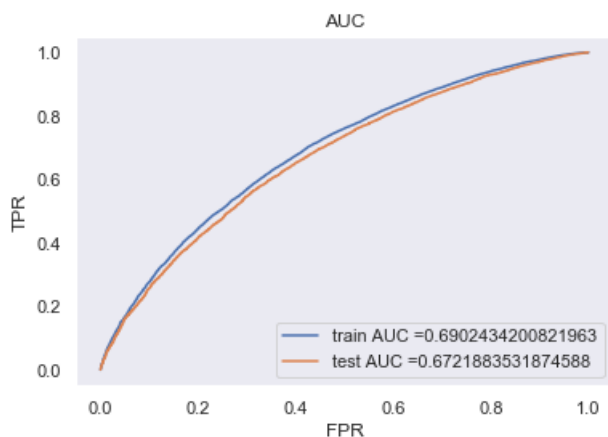
```
def pred_prob(clf, data):
    y_pred = []
    y_pred = clf.predict_proba(data)[:,1]
    return y_pred

def find_best_threshold(threshold, fpr, tpr):
    t = threshold[np.argmax(tpr*(1-fpr))]
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.round(t,3))
    return t

def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        if i>=threshold:
            predictions.append(1)
        else:
            predictions.append(0)
    return predictions
```

In [89]:

```
#
https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc\_curve.html#sklearn.metrics.roc\_curve
from sklearn.metrics import roc_curve, auc
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(max_depth = max_d, n_estimators = n_est)
model.fit(X_train,y_train)
y_train_pred = pred_prob(model,X_train)
y_test_pred = pred_prob(model,X_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.close()
plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("AUC")
plt.grid()
plt.show()
```



In [168]:

```
#Train Confusion Matrix of TFIDE
best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
```


the maximum value of $tpr \cdot (1 - fpr)$ 0.40775354295248856 for threshold 0.848

In [161]:

```
print("Train confusion matrix")
co=confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t))
sns.heatmap(co, annot=True, annot_kws={"size": 25}, fmt="d", linewidths=.5, yticklabels=2)
```

Out[161]:

<matplotlib.axes._subplots.AxesSubplot at 0x7212ec6a48>




In [162]:

```
#Test Confusion Matrix of TFIDF
print("Test confusion matrix")
sns.heatmap(confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t)), annot=True, annot_kws={"size": 25}, fmt="d", linewidths=.5, yticklabels=2)
```

Test confusion matrix

Out[162]:

<matplotlib.axes._subplots.AxesSubplot at 0x7212f918c8>



In [163]:

```
train_tfidf_w2v_essays_np = np.array(train_tfidf_w2v_essays)
test_tfidf_w2v_essays_np = np.array(test_tfidf_w2v_essays)
```

In [196]:

```
from scipy.sparse import coo_matrix, hstack
tr1 = coo_matrix(X_train["cat_0"].values.reshape(-1,1))
tr2 = coo_matrix(X_train["cat_1"].values.reshape(-1,1))
tr3 = coo_matrix(X_train["subcat_0"].values.reshape(-1,1))
tr4 = coo_matrix(X_train["subcat_1"].values.reshape(-1,1))
tr5 = coo_matrix(X_train["state_0"].values.reshape(-1,1))
```



```

iid='warn', n_iter=10, n_jobs=-1,
param_distributions={'max_depth': [2, 5, 10],
                    'n_estimators': [10, 50, 100]},
pre_dispatch='2*n_jobs', random_state=None, refit=True,
return_train_score=True, scoring='roc_auc', verbose=0)

```

In [201]:

```

print('Best score: ',rs.best_score_)
print('k value with best score: ',rs.best_params_)
print('='*75)
print('Train AUC scores')
print(rs.cv_results_['mean_train_score'])
print('CV AUC scores')
print(rs.cv_results_['mean_test_score'])

```

```

Best score: 0.6174958760137832
k value with best score: {'n_estimators': 50, 'max_depth': 2}
=====
Train AUC scores
[0.6171561 0.65694083 0.68747866 0.68388587 0.83754251 0.92801416
 0.94418432 0.99999969 1.          ]
CV AUC scores
[0.607904 0.61749588 0.61233596 0.60930948 0.59282055 0.57987553
 0.57181274 0.55986971 0.56357041]


```

In [202]:

```

import seaborn as sns; sns.set()
max_scores1 = pd.DataFrame(rs.cv_results_).groupby(['param_n_estimators', 'param_max_depth']).max()
max_scores1.unstack(['mean_test_score', 'mean_train_score'])
fig, ax = plt.subplots(1,2, figsize=(20,6))
sns.heatmap(max_scores1.mean_train_score, annot = True, fmt='.4g', ax=ax[0])
sns.heatmap(max_scores1.mean_test_score, annot = True, fmt='.4g', ax=ax[1])
ax[0].set_title('Train Set')
ax[1].set_title('CV Set')
plt.show()

```



In [203]:

```

#Best Parameters for TFIDF-W2V
rs.best_params_

```

Out[203]:

```
{'n_estimators': 50, 'max_depth': 2}
```

In [204]:

```


max_d = rs.best_params_['max_depth']
n_est = rs.best_params_['n_estimators']

```

In [205]:

In [205]:

```
#
https://scikitlearn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
model = RandomForestClassifier(max_depth = max_d, n_estimators = n_est)
model.fit(X_train,y_train)
y_train_pred = pred_prob(model,X_train)
y_test_pred = pred_prob(model,X_test)
train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
plt.close
plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("AUC")
plt.grid()
plt.show()
```




In [206]:

```
#Train TFIDF-W2V Confusion Matrix
best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
print("Train confusion matrix")
co=confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t))
sns.heatmap(co, annot=True, annot_kws={"size": 25},fmt="d",linewidths=.5,yticklabels=2)
```

the maximum value of $tpr * (1 - fpr)$ 0.32673833475136643 for threshold 0.848
Train confusion matrix

Out[206]:

<matplotlib.axes._subplots.AxesSubplot at 0x7230868ac8>



In [207]:

```
#Test TFIDF-W2V Confusion Matrix
print("Test confusion matrix")
```

```
print(test_confusion_matrix)
sns.heatmap(confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t)), annot=True, annot_kws={"size": 25}, fmt="d", linewidths=.5, yticklabels=2)
```

Test confusion matrix

Out[207]:

<matplotlib.axes._subplots.AxesSubplot at 0x723082e548>



3. Summary

In [1]:

```
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Vectorizer", "Model", "Hyperparameters(n_estimators,max_depth)", "Test AUC"]
x.add_row(["TFIDF", "GBDT", "(100, 10)", 0.672])
x.add_row(["TFIDF W2V", "GBDT", "(100, 5)", 0.605])
print(x)
```

Vectorizer	Model	Hyperparameters(n_estimators,max_depth)	Test AUC
TFIDF	GBDT	(100, 10)	0.672
TFIDF W2V	GBDT	(100, 5)	0.605

In []: